

SHASHWAT YADAV

Mobile : +91 6393355490
Email : shashwatyadav101@gmail.com

Linked : shashwatyadav101
GitHub : Shashwat-911

SUMMARY

AI/ML Engineering student (B.E. 2027, CGPA 9.15/10) with hands-on production experience in backend engineering, distributed systems, and ML integration. Proficient in Python, Java, FastAPI, Django, and microservices architecture. Built production-grade systems including a clustered caching engine, sharded vector search, and enterprise LLM pipelines. IEEE-published, two-time hackathon winner, with strong fundamentals in DSA, system design, and databases.

SKILLS SUMMARY

- AI ML:** Supervised Learning, Deep Learning, NLP, RAG Pipelines, LangChain, LLMs, Prompt Engineering, Semantic Search, VectorDB
- Languages:** Python (primary), Java, SQL, JavaScript, C
- Frameworks:** Django, FastAPI, REST APIs, gRPC, Microservices, Node.js
- Databases:** MongoDB, Vector DB, MySQL
- Soft Skills:** Time Management, Problem Solving, Analytical Thinking, Team Management, Leadership, Adaptability

EXPERIENCE

AI Engineer Intern — Backend & ML Systems

Jan 2026 – May 2026

Hasprana Health Care Solutions · Bengaluru, India

- Designed & deployed ML models for predictive analytics in production — improved forecasting accuracy by 15% and operational efficiency by 20%.
- Built Python & Java automation pipelines for healthcare data processing; cut manual effort by 30%, enabling real-time analytics on 10,000+ daily records.
- Integrated LLM modules into enterprise apps within a 6–8 person Agile/Scrum squad, delivering 100% of sprint goals across a 5-month engagement.
- Authored 40+ unit & integration tests achieving 95%+ code coverage on critical data-pipeline modules.

PROJECTS

- Medical RAG System** | [Github](#)
Tech : Python · LangChain · ChromaDB · Ollama · Groq API
 - Built a production-grade RAG pipeline (ingestion → retrieval → LLM inference) grounded in clinical data, achieving near-zero hallucinations across 500+ queries.
 - Improved efficiency via optimized chunking, embeddings, and prompting, reducing latency by ~40% while boosting reliability..
- AeroCache — Clustered In-Memory Caching System** | [Github](#)
Tech : Java · Consistent Hashing · LRU/LFU Eviction · Multi-Node Systems
 - Designed a distributed cache with consistent hashing for seamless scaling and zero-downtime node changes.**
 - Achieved 99.9% uptime across 10 nodes with efficient eviction strategies and multi-node coordination..**
- NexusSearch — Sharded Vector Search Engine** | [Github](#)
Tech : Python · gRPC · ANN · Horizontal Sharding
 - Built a distributed ANN-based vector search engine with sharding, reducing latency by ~35% and scaling to 100K+ records.
 - Developed fault-tolerant gRPC microservices enabling sub-50ms query performance.
- Sandboxed Algorithmic Judge & AI Code Reviewer** | [Github](#)
Tech : Python · Groq/Llama · Docker · REST API
 - Created an LLM-powered automated code review system evaluating correctness, complexity, and code quality across multiple languages using structured generative analysis.
 - Built a secure sandboxed execution environment running 50+ test cases per submission in under 2 seconds, ensuring isolation and zero cross-submission data leakage. .

EDUCATION

- | | |
|--|-----------------------|
| B.E. — Artificial Intelligence & Machine Learning CGPA: 9.15 | Bengaluru |
| Dayananda Sagar Academy of Technology & Management, Bengaluru | Oct 2023 – June 2027* |
| Favourite Coursework: Machine Learning, Deep Learning, NLP, DSA, Linear Algebra, Probability & Statistics, DBMS, OS, System Design | |
| Kendriya Vidyalaya | Bengaluru |
| 12th Standard — Physics, Chemistry, Mathematics (CBSE) Percentage: 78% | April 2021–March 2023 |
| Kendriya Vidyalaya | Bengaluru |
| 10th Standard (CBSE) Percentage: 78% | April 2020–March 2021 |

ACHIEVEMENTS & PUBLICATIONS

- BMS Hackathon — 1st Place:** Medical Report Simplifier (LLM + Django); outperformed all 20+ competing teams.
- IEEE Publication — NQComp 2026:** "A Review on AI-Enabled Wildlife Preservation and Management System" — IEEE International Conference on Next-Gen Quantum and Advanced Computing.
- Smart India Hackathon 2025 — Institutional Nominee:** Selected from 100+ teams for innovative AI/ML technical solution.

CERTIFICATIONS

- Certified System Administrator — ServiceNow
- Agentic AI Executive — ServiceNow
- Machine Learning Fundamentals — Infosys Springboard
- Introduction to NLP — Infosys Springboard